

PDF Classifier using Trap4Phish Dataset

Shashank Mishra

Shubham Negi

Christopher Nyandoro

Table of Contents



Objective



Data Collection



Exploratory Data Analysis



Data Cleaning & Processing



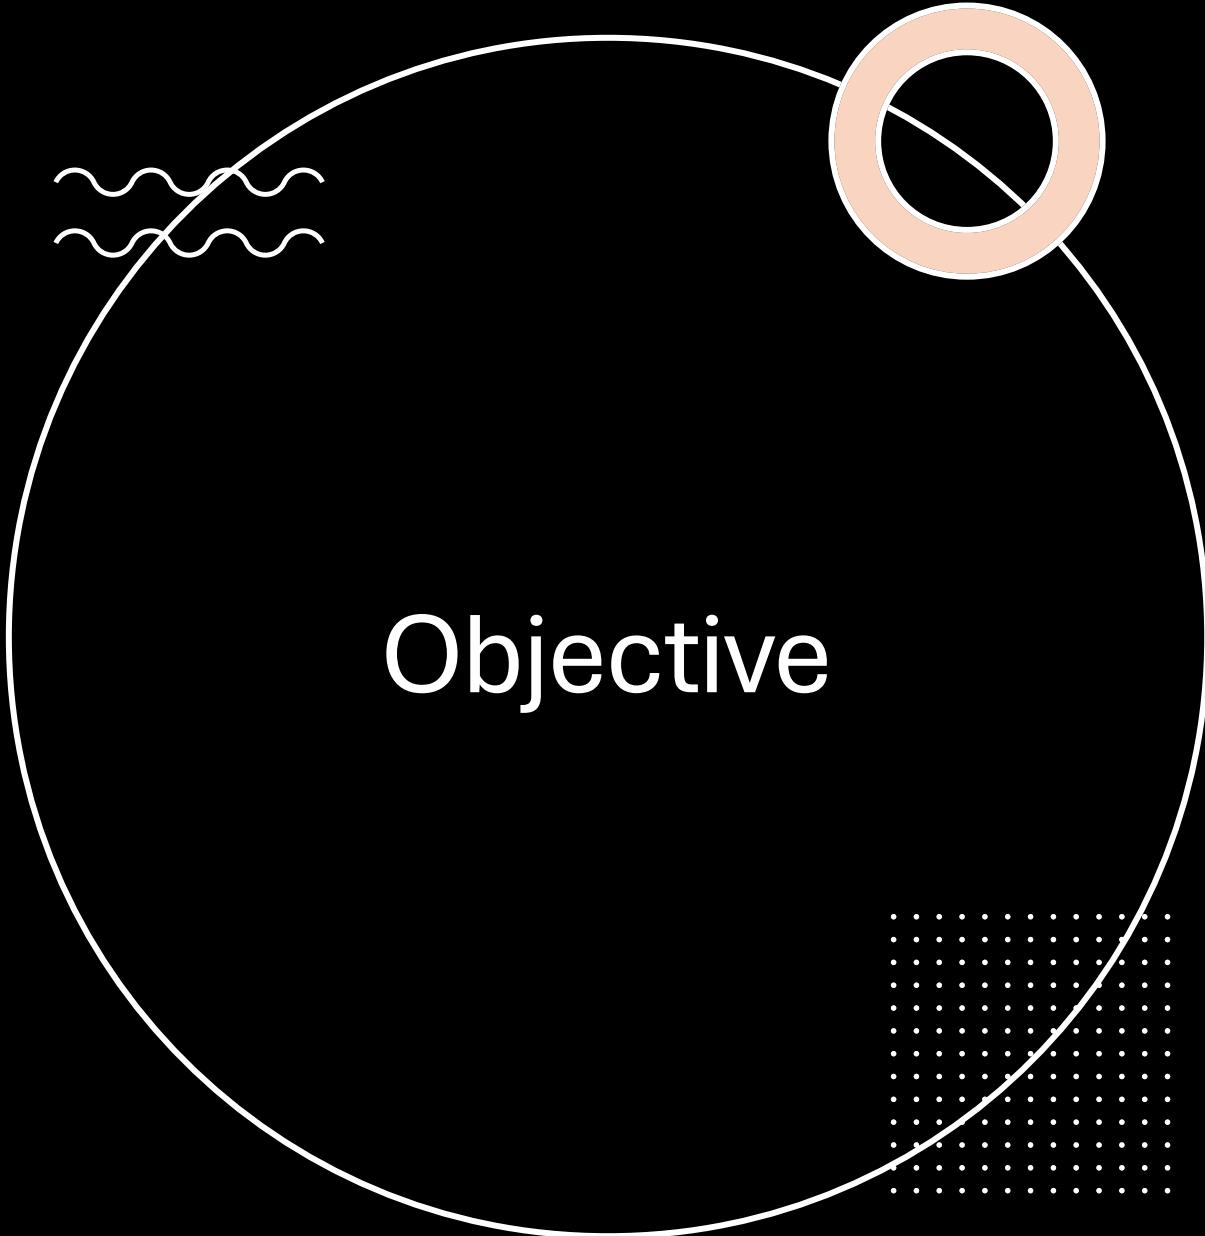
Model Selection & Training



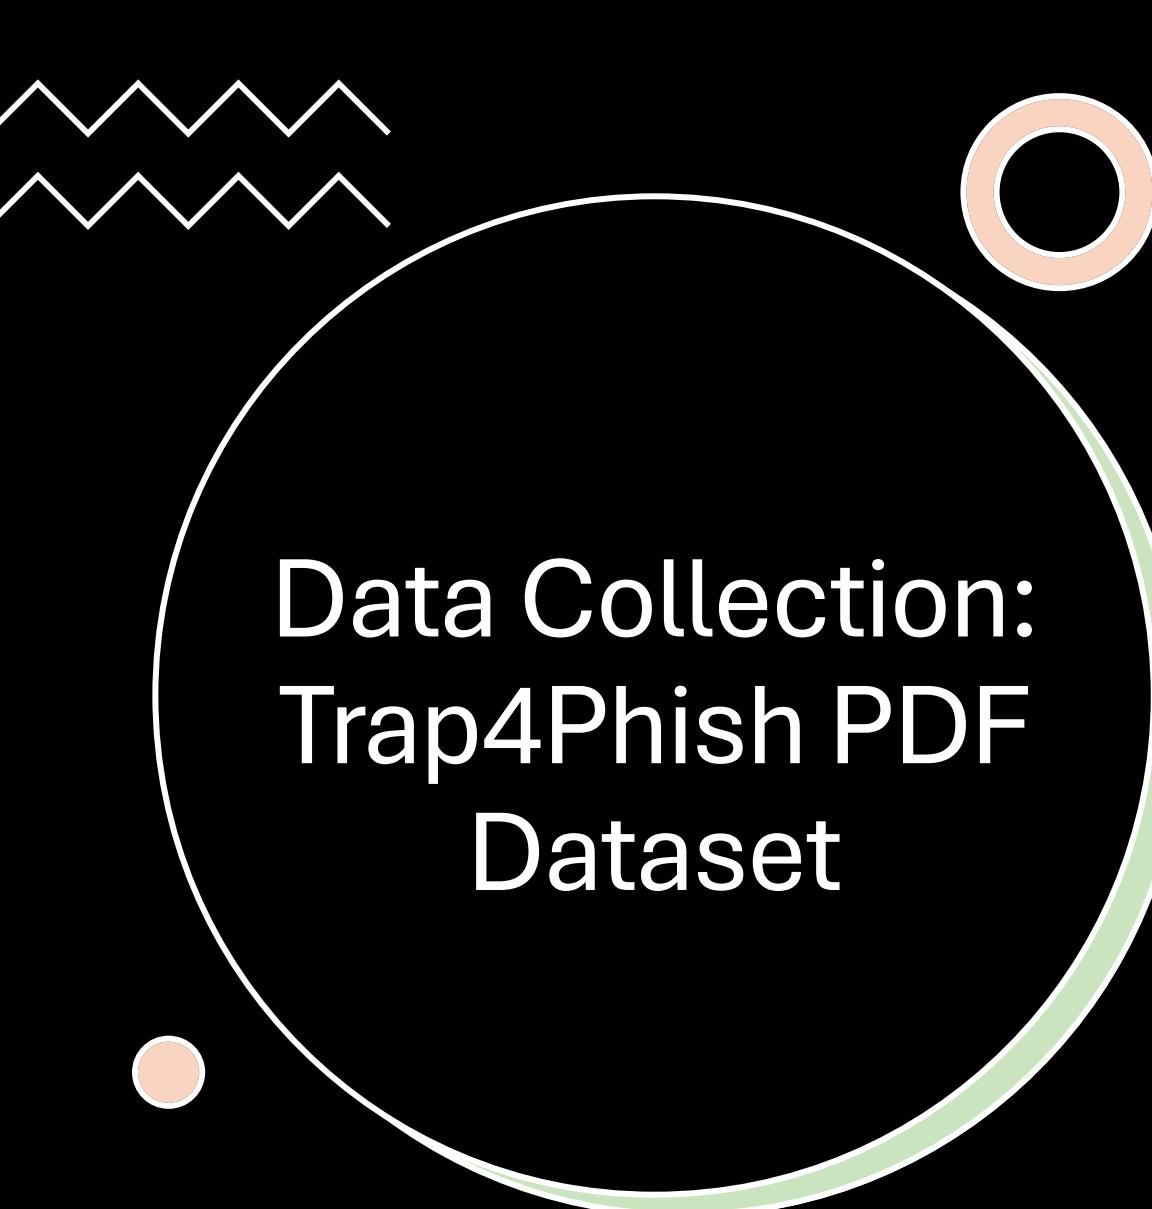
Evaluation



ML Pipeline



Build an end-to-end ML pipeline which provided a PDF, can extract features from the PDF, give it as input to the ML Model, and then classify the PDF as malicious or benign, and also provide a confidence score, with a good recall score.



Data Collection: Trap4Phish PDF Dataset

- Contains 19296 rows of data
- Contains 42 features
- Benign PDF Source: Legitimate files collected from trusted public and institutional repositories (government data portals, academic datasets, open data sources).
- Malicious PDF Source: Files obtained from verified malware and phishing feeds such as MalwareBazaar, PhishTank, and VirusShare





Data Collection: Trap4Phish PDF Dataset

Features range from Structural Metadata, Content Objects, Behavioral Keywords, and Known Malicious Indicators

Examples

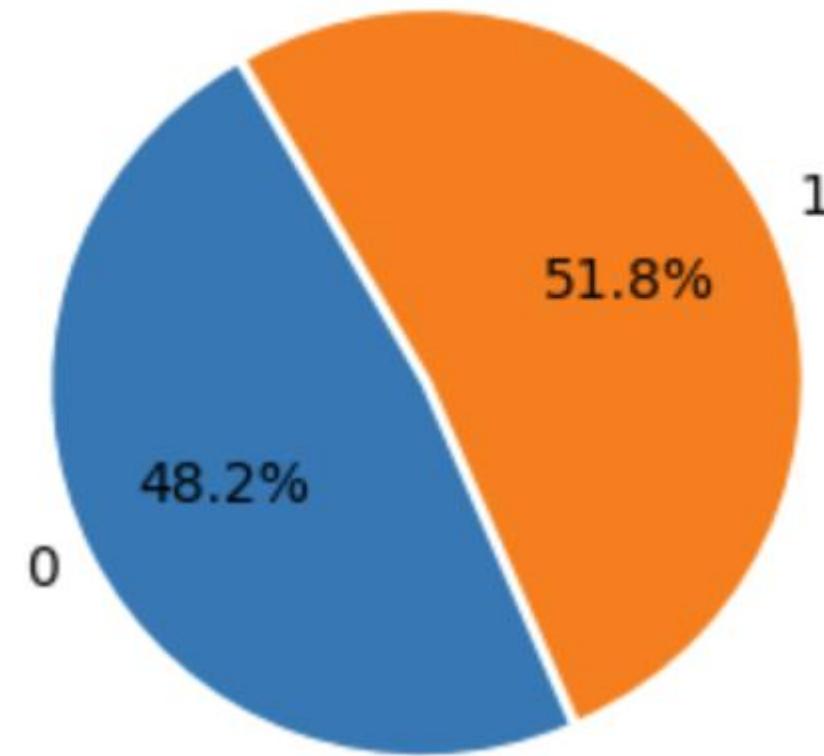
- Structural Metadata: file_size, page_count, title_chars, metadata_size
- Content Objects: object_count, stream_count, image_count
- Behavioral Keywords: javascript_count, openaction_count, aa_count
- Malicious Indicators: name_obfuscations, jbig2decode_count



Exploratory Data Analysis (EDA)

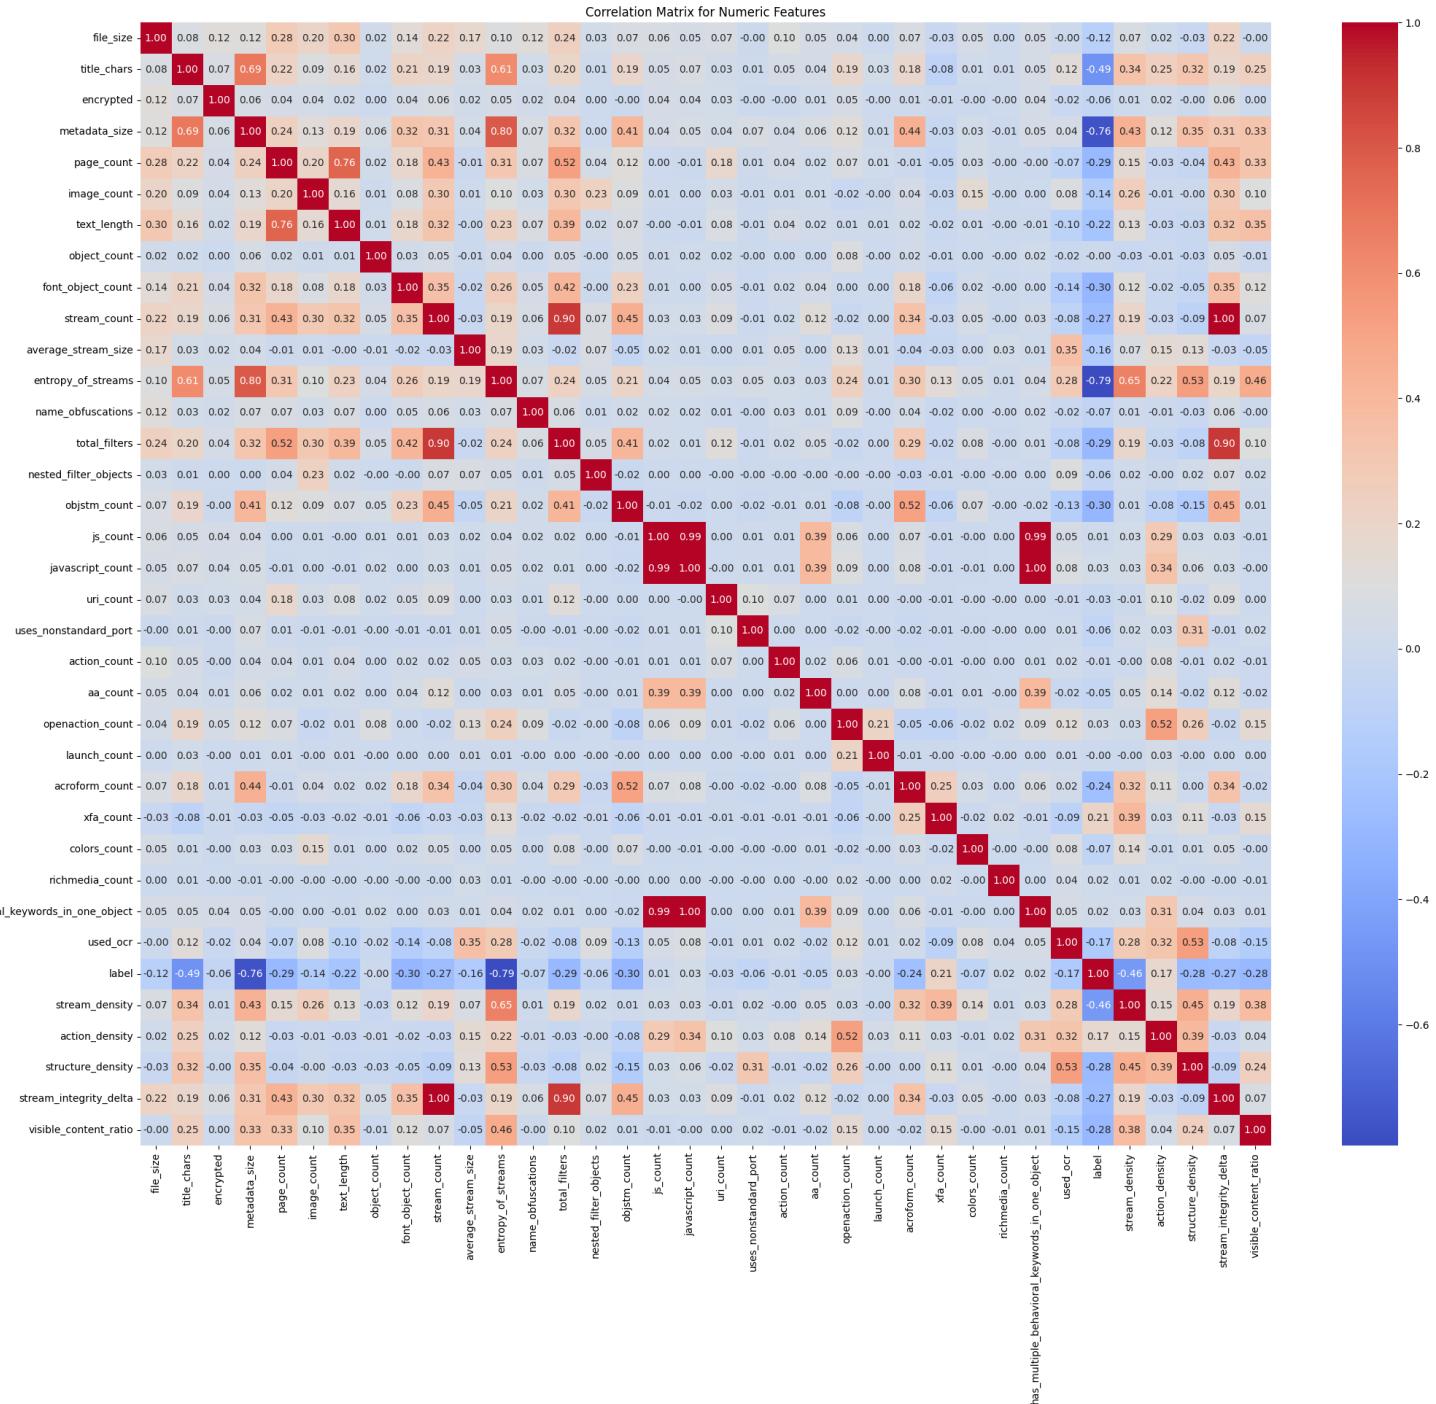
- Class Distribution
- Feature Correlations
- `ydata.profiling`
- Box Plot for label
- Feature Importance using ML Classifier
- Scatter Plot
- Investigating the feature
``used_ocr``

Dataset Balance: Benign vs. Malicious



EDA: Class
Distributions

EDA: Feature Correlations



EDA: ydata.profiling

Dataset statistics

Number of variables	37
Number of observations	19296
Missing cells	0
Missing cells (%)	0.0%
Total size in memory	5.4 MiB
Average record size in memory	296.0 B

Variable types

Numeric	37
----------------	----

EDA: ydata.profiling – Problem: Rare Signals

encrypted

Real number (\mathbb{R})

Zeros

Distinct 2

Distinct (%) < 0.1%

Missing 0

Missing (%) 0.0%

Infinite 0

Infinite (%) 0.0%

Mean 0.004353233831

Minimum 0

Maximum 1

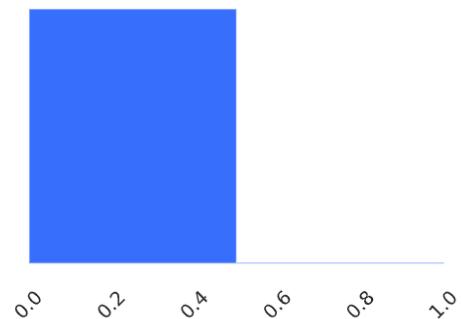
Zeros 19212

Zeros (%) 99.6%

Negative 0

Negative (%) 0.0%

Memory size 150.9 KiB



More details

Statistics

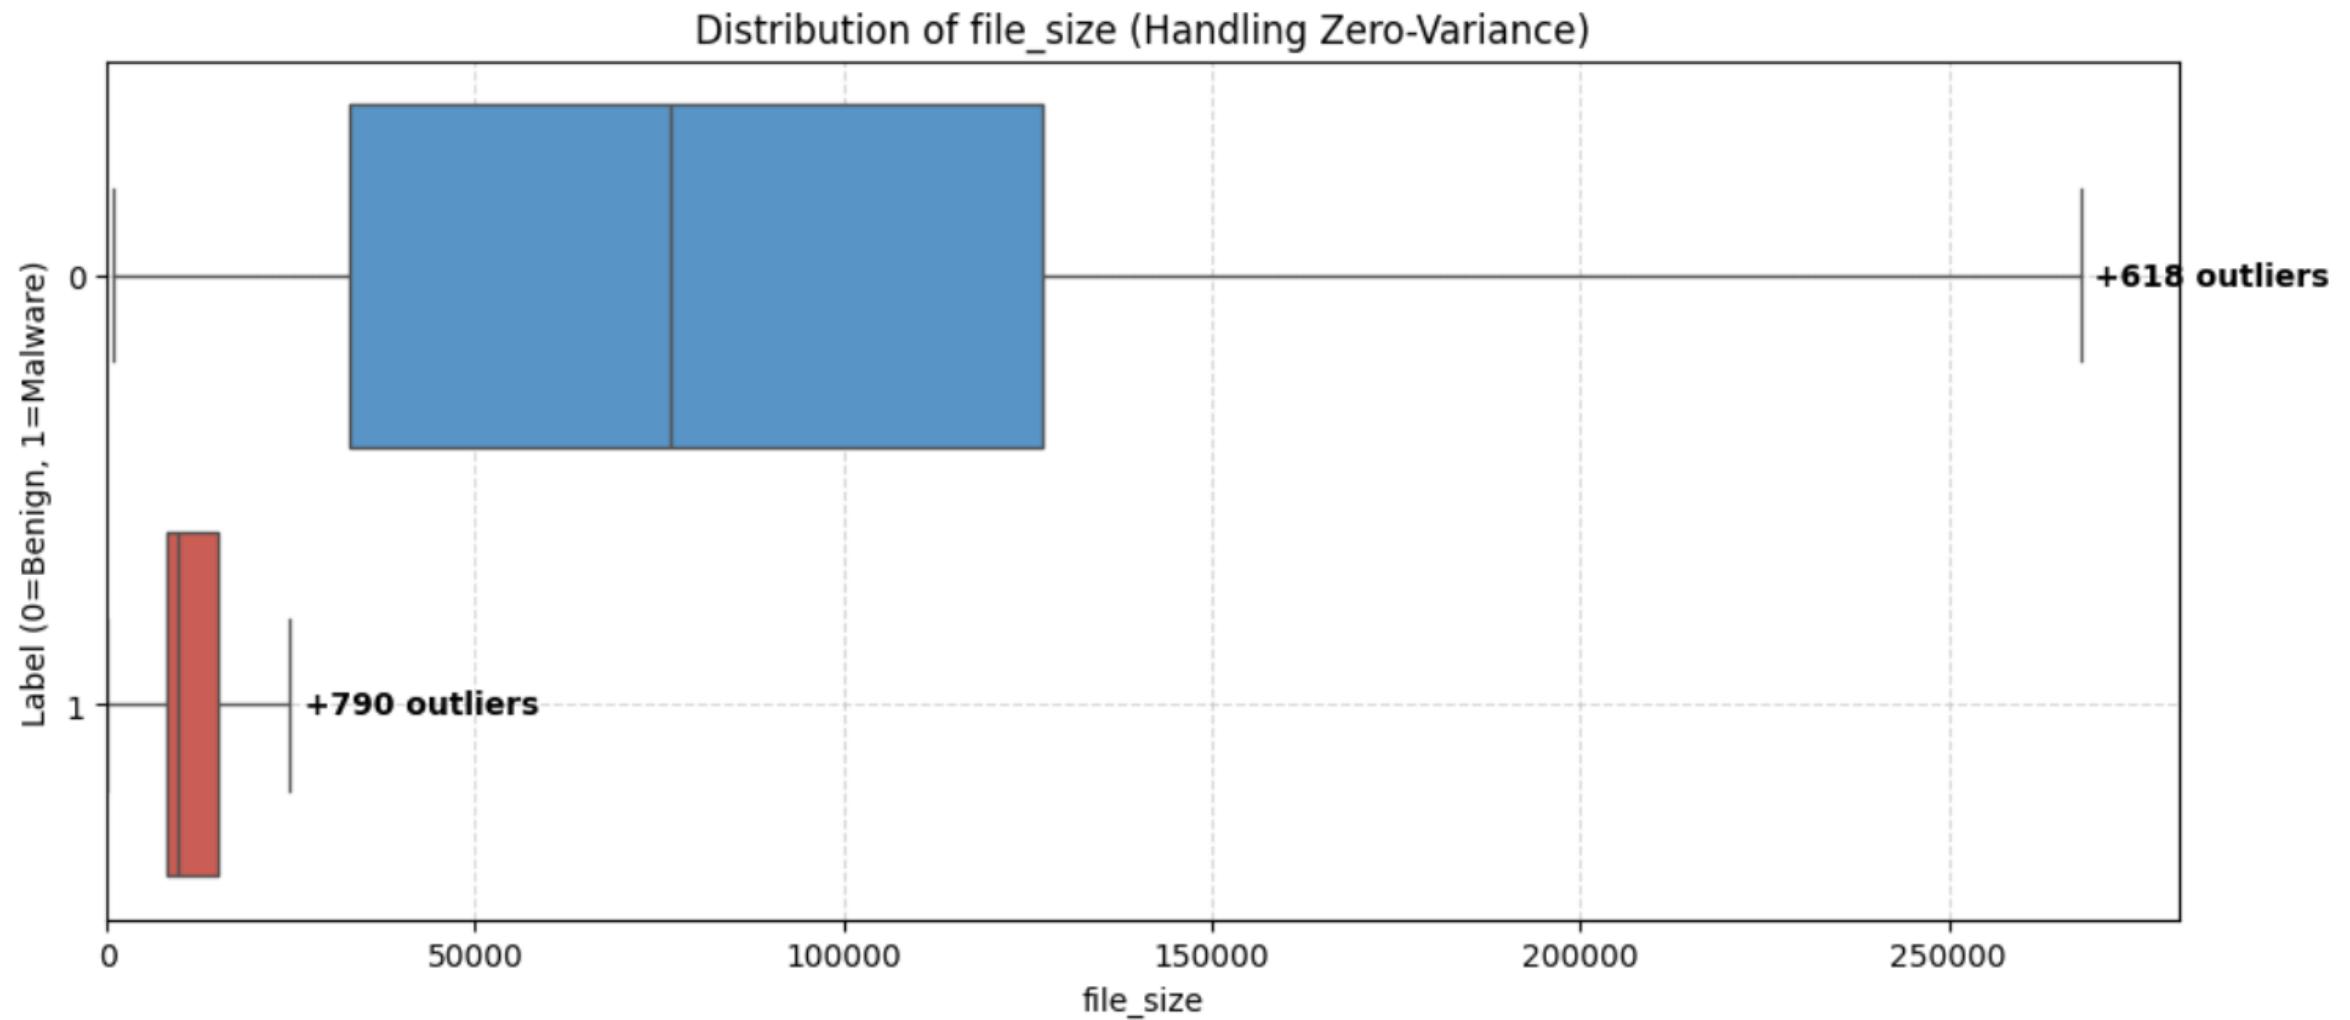
Histogram

Common values

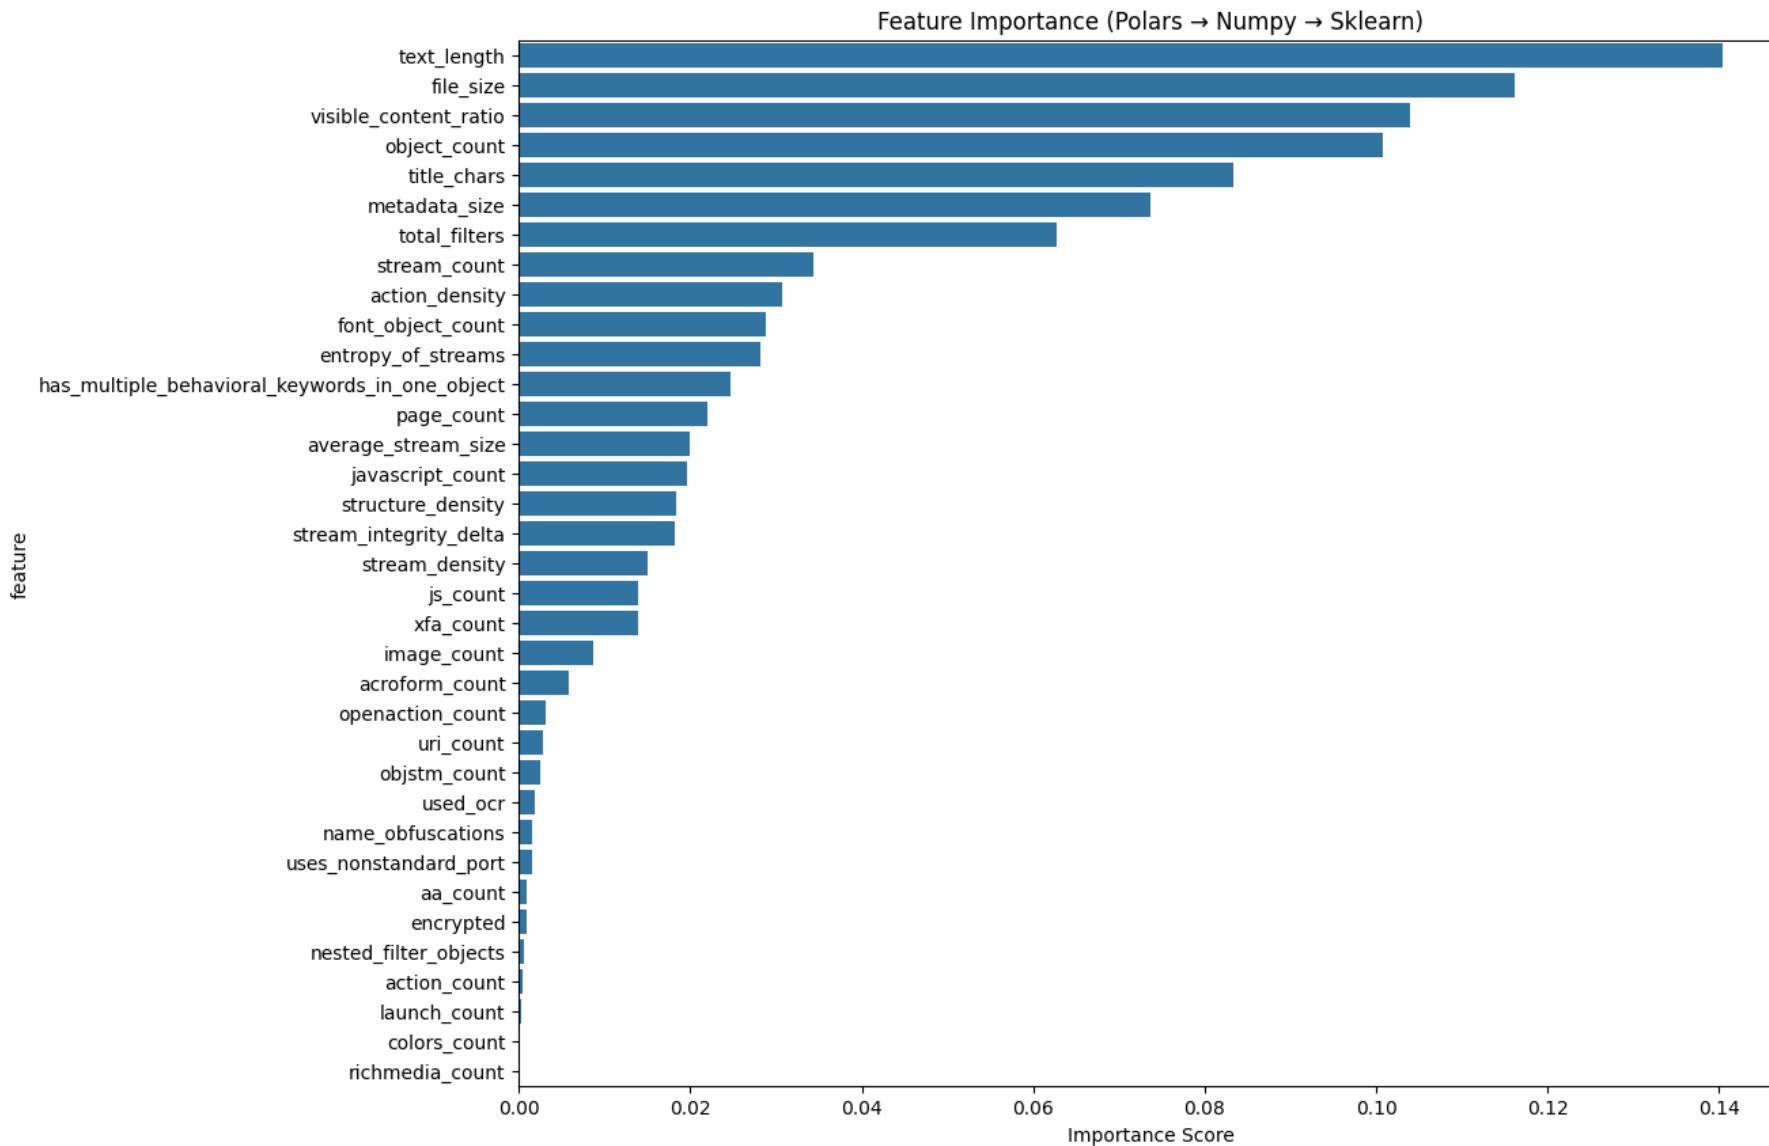
Extreme values

Value	Count	Frequency (%)
0	19212	99.6%
1	84	0.4%

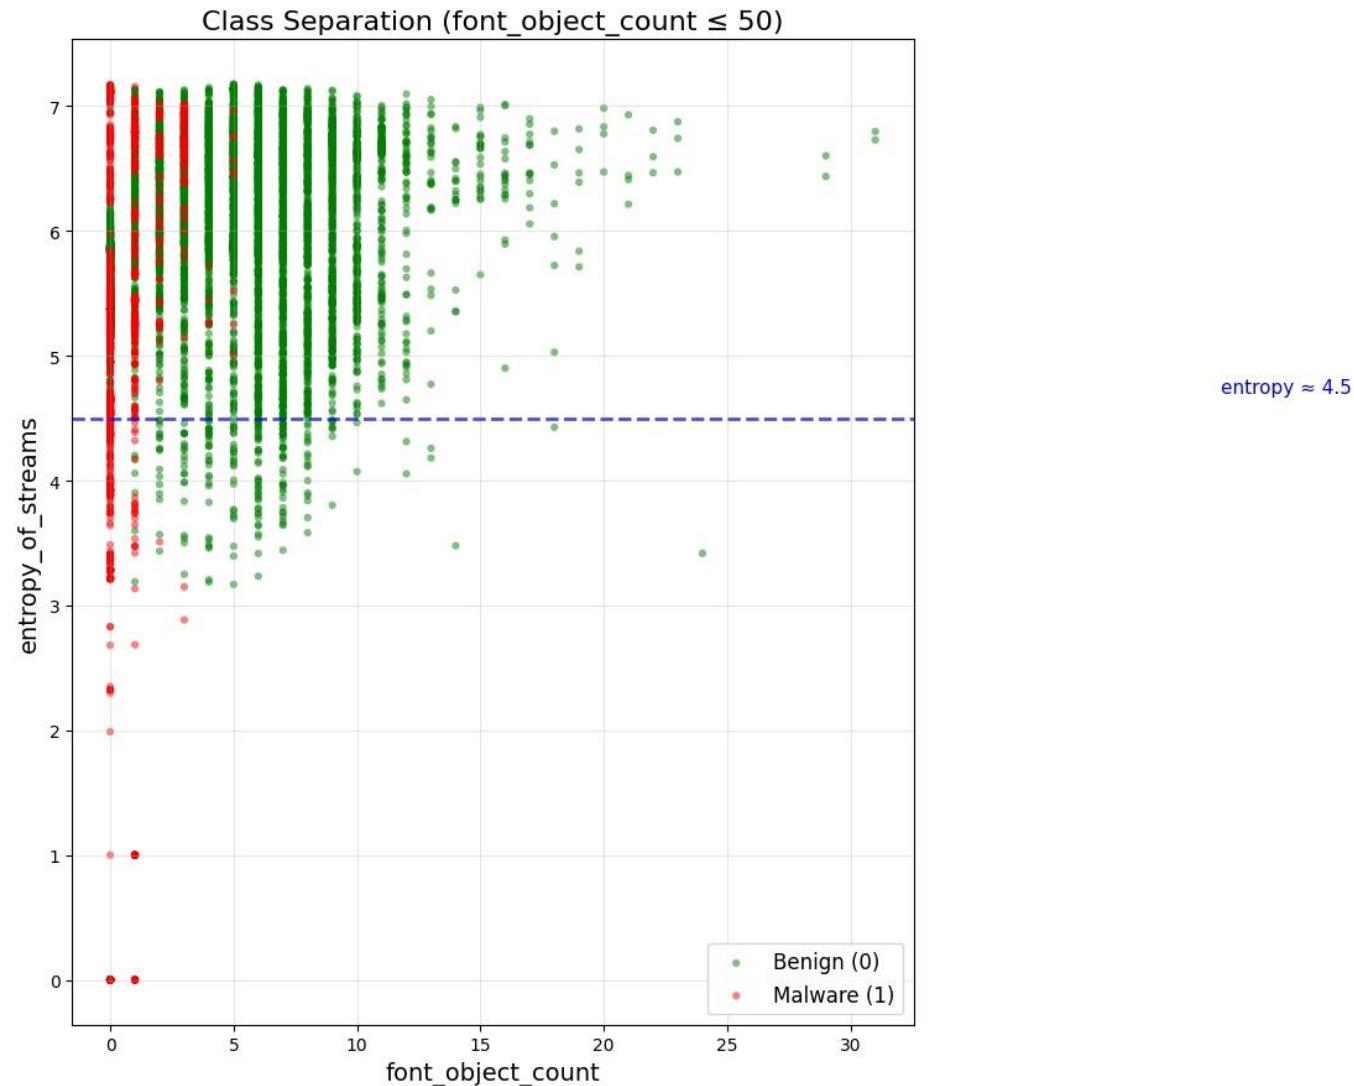
EDA: Box Plot



EDA: Feature Importance using ML



EDA: Scatterplot – Clear Separation problem



EDA: Investigating the feature `used_ocr`

- Binary feature which tells if OCR used to extract features from the PDF.
- Cause of concern: Exception statement in Extraction Code

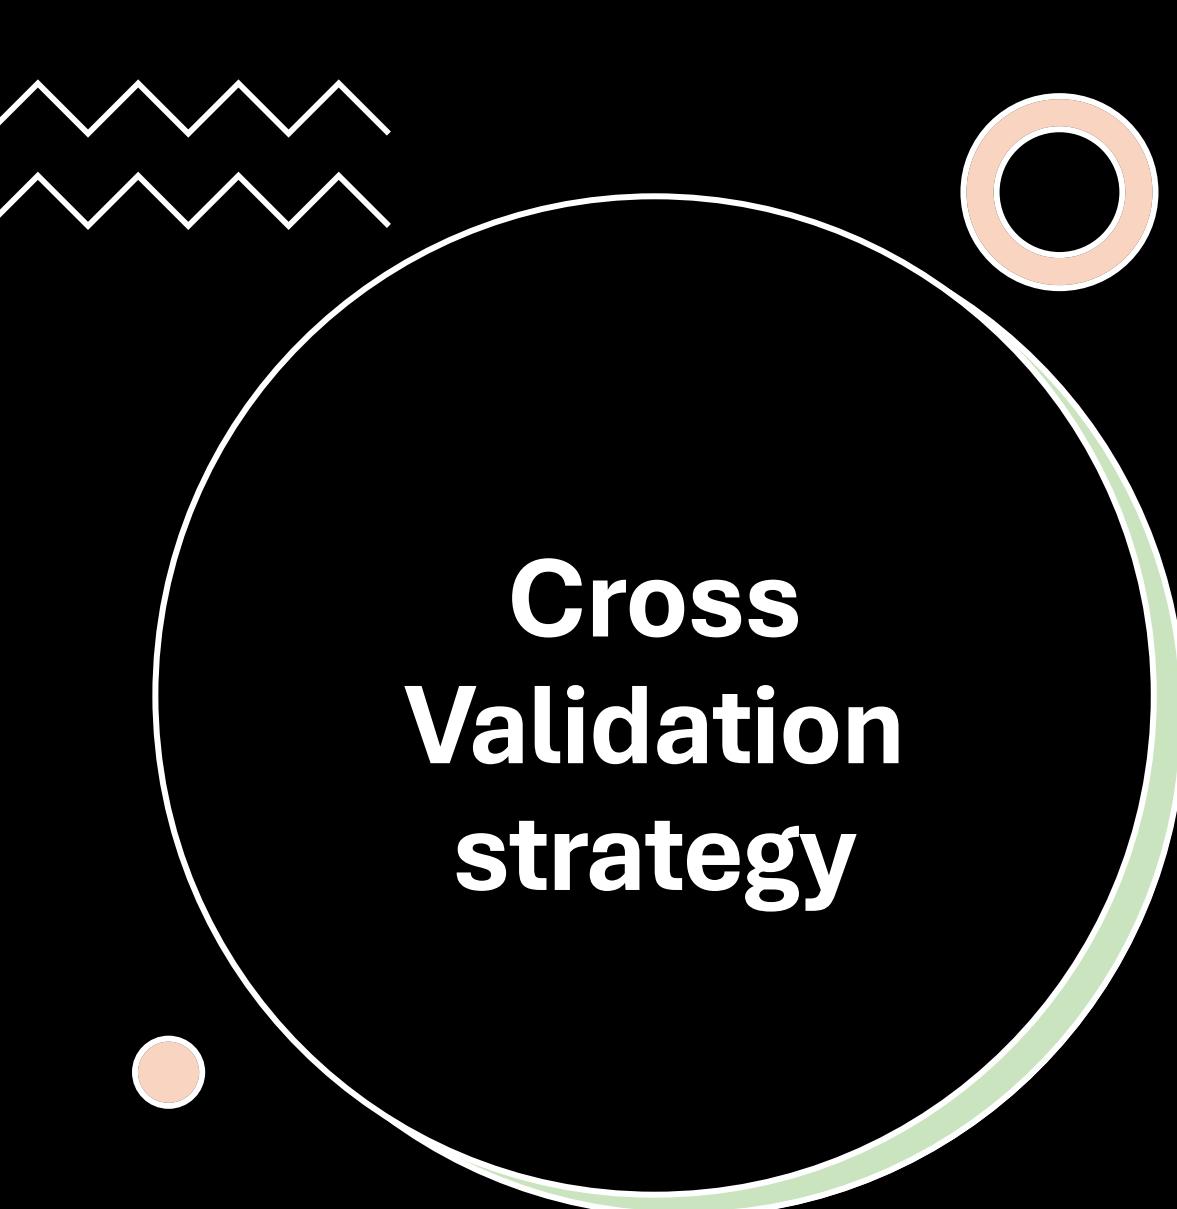
[Refer to ocr.ipynb]



Data Cleaning & Processing

- 1. Missing values -> Removed Empty Columns
- 2. Duplicates -> Duplicates removed
- 3. Inconsistencies -> used_ocr investigated
- 4. Transformation -> log1p used to fix right-skeweness
- 5. Outlier Management -> Used IQR to Winsorize

[Refer to Model3.ipynb]



Cross Validation strategy

- **Strategy:** StratifiedKFold(`n_splits=10, shuffle=True, random_state=42`) preserves class ratios per fold.
- **Metrics tracked:** Recall (priority), Precision, F1, Accuracy, ROC-AUC.
- **Rationale:** Stratification mitigates imbalance bias; 5-fold balances variance and runtime.
- **Robustness checks:** Tried 10-fold and repeated stratified CV; results consistent.



Model Selection & Training

MODEL SELECTION & HYPERPARAMETER TUNING ANALYSIS

PART 1: COMPARING DIFFERENT ALGORITHMS

Evaluating models with 5-fold Stratified Cross-Validation...

Model	Mean Acc	Std	Time (s)
Logistic Regression	0.9702	0.0044	7.88
Decision Tree	0.9962	0.0012	0.34
Random Forest	0.9975	0.0007	4.41
Gradient Boosting	0.9969	0.0007	10.44
K-Nearest Neighbors	0.9816	0.0016	0.99
Naive Bayes	0.7023	0.0102	0.15

- **Feature Standardization** (z-score normalization)
- **Stronger Regularization** (deeper pruning, higher min_samples)
- **Simpler Model** (fewer trees, shallower depth)
- **Lower decision threshold** for malware detection

5 KEY INSIGHTS

- ✓ Tree-based models (RF, GB) significantly outperform linear models
- ✓ Hyperparameter tuning provides marginal improvement (~0.1-0.3%)
- ✓ High accuracy is consistent across different CV strategies
- ✓ Model generalizes well (minimal overfitting)
- ✓ All evaluation metrics exceed 99% - excellent performance

Hyperparameters Tuned

- `n_estimators`: More trees improve stability; chosen in the 100–200 range.
- `max_depth`: Cap to prevent memorization; chosen around 10–15 based on CV.
- `min_samples_leaf`: Small but >1 (e.g., 3–5) to reduce variance and catch minority patterns.
- `min_samples_split`: 5–10 to avoid overly fine splits.
- `max_features`: '`sqrt`' for decorrelated trees, improving generalization.
- `class_weight`: `{0: 1, 1: 3}` to penalize malware misclassification; raises recall with acceptable precision loss.

Optuna Tuning

- **Objective:** Maximize CV recall on class 1 (malware).
- **Sampler/Pruner:** `TPE` sampler with `MedianPruner` for early stopping of weak trials.
- **Search space:** `n_estimators` (50–200), `max_depth` (5–20 or `None`), `min_samples_leaf` (1–10), `min_samples_split` (2–20),
`max_features` (`'sqrt'`, `'log2'`, 0.5), `class_weight` (e.g., `{0:1,1:2..5}`).
- **Outcome:** Efficient exploration converging on recall-focused configs; similar/better recall than grid/random with fewer trials.

Class Weighting & Threshold

- **Class weighting:** Prefer `{0:1, 1:3}` as balanced point; higher weights (≥ 5) can overinflate FPs.
- **Threshold tuning:** Move decision threshold below 0.5 (e.g., ≈ 0.15) to reduce false negatives; select via precision-recall curve.
- **Process:** Sweep thresholds, report FN/FP/Recall/Precision; pick threshold minimizing FN with acceptable FP.

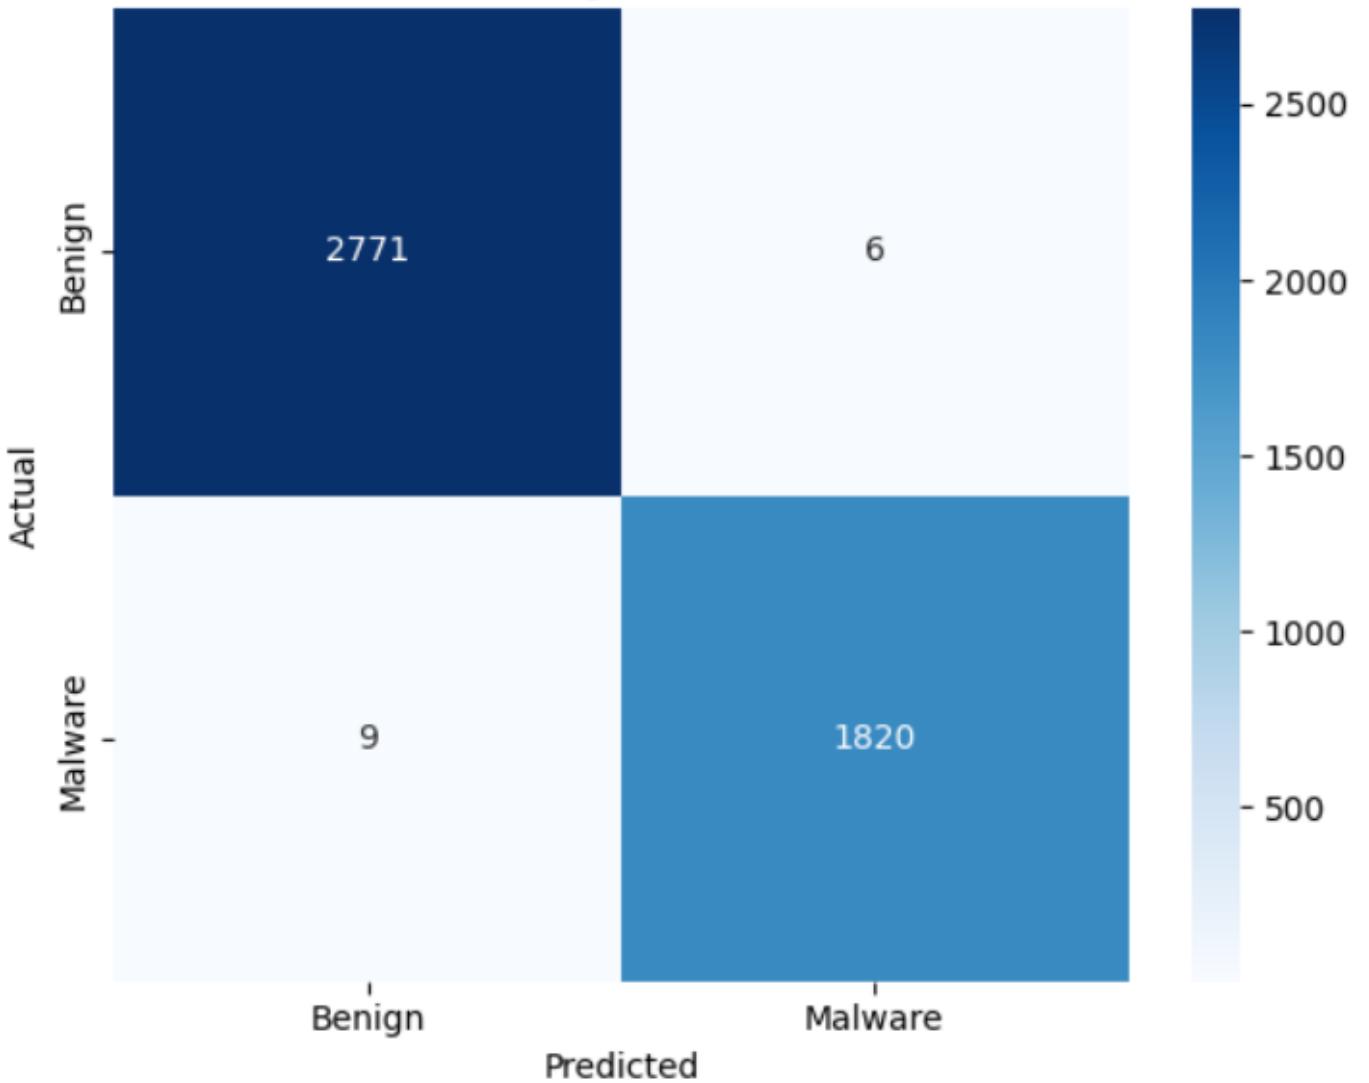


Evaluation

- **Primary metric: Recall:** Optimize for catching malware; minimize false negatives over raw accuracy.
- **Cost model:** False negative is high-risk (missed malware); false positive is lower-cost (extra review).
- **Trade-off:** Accept more false positives if it materially reduces false negatives.

ROC-AUC Score: 0.9991

GradientBoosting - Confusion Matrix





No Leakage Detected

Check	Result
Features with $ r > 0.95$	None found
Single dominant feature (>50% importance)	No — top feature has 16%
Train-test gap	0.06% (minimal)

Ablation Study Results

Features Removed	Accuracy
0 (baseline)	99.74%
Top 5	99.64%
Top 10	99.65%
Top 15	99.09%
Top 20	80.82%

A large industrial piping system, painted blue, stands against a backdrop of a clear blue sky and a setting or rising sun. The pipes are interconnected with various fittings, valves, and support structures, creating a complex network. The lighting highlights the metallic surfaces of the pipes.

Production Pipeline Implementation

Model Import & Prediction
System

Production Pipeline Architecture

```
• src/
  • └── features.py          # Extract 41 features from PDF
  • └── preprocessing.py     # Transform pipeline (Winsorize → log1p → scale)
  • └── import_model.py      # Import trained models
  • └── predict.py           # Prediction CLI

  • models/                  # Saved artifacts
  • └── model.joblib          # The trained model
  • └── caps.json             # Transformation parameters
  • └── scaler.joblib         # Fitted StandardScaler
  • └── features.json         # Feature names
```

Core Design

- **Goal:** Import any trained model and ensure consistent preprocessing
- **Challenge:** Training preprocessing must match prediction preprocessing exactly
- **Solution:** Save all transformation parameters alongside model

preprocessing.py

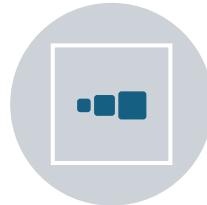
- The Pipeline

```
• class PDFPreprocessingPipeline:  
•     def fit(self, df):  
•         # Training mode: Learn from data  
•         self.transformer.fit(df)          # Learn 99th percentile caps  
•         self.scaler = StandardScaler()   # Fit scaler  
  
•     def transform(self, features):  
•         # Prediction mode: Apply saved transformations  
•         # 1. Winsorize + log1p (using saved caps)  
•         # 2. Select 28 features  
•         # 3. Scale (using saved scaler)  
•         return X_scaled    # numpy array → model
```

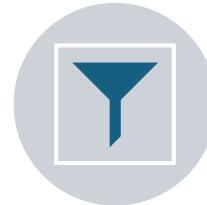
Key Point: Same pipeline code for training and prediction, just different mode (fit vs transform)

import_model.py - Model Import

```
python src/import_model.py model.joblib --test
```



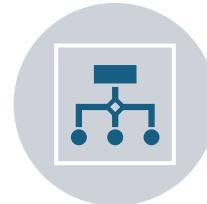
1. Load the trained model



2. Fit pipeline on training dataset → learns caps & scaler



3. Test compatibility (feature count, predictions)



4. Save everything to models/:
 - model.joblib
 - caps.json (transformation parameters)
 - scaler.joblib (fitted scaler)
 - features.json (feature names)

```
Sample predictions:
```

```
INFO - ✓ Sample 0: True=Benign, Predicted=Benign
INFO - ✓ Sample 1: True=Benign, Predicted=Benign
INFO - ✓ Sample 2: True=Benign, Predicted=Benign
INFO - ✓ Sample 3: True=Benign, Predicted=Benign
INFO - ✓ Sample 4: True=Malicious, Predicted=Malicious
INFO - ✓ Sample 5: True=Benign, Predicted=Benign
INFO - ✓ Sample 6: True=Malicious, Predicted=Malicious
INFO - ✓ Sample 7: True=Benign, Predicted=Benign
INFO - ✓ Sample 8: True=Malicious, Predicted=Malicious
INFO - ✓ Sample 9: True=Malicious, Predicted=Malicious
```

```
INFO -
```

```
Sample accuracy: 100.0%
```

```
INFO - =====
```

```
INFO -
```

```
Saving model to models...
```

```
INFO - ✓ Model saved: models/model.joblib
INFO - ✓ Features saved: models/features.json
INFO - ✓ Transformation caps saved: models/caps.json
INFO - ✓ Scaler saved: models/scaler.joblib
INFO - ✓ Metadata saved: models/metadata.json
```

```
INFO -
```

```
=====
```

```
INFO - IMPORT COMPLETE!
```

```
INFO - =====
```

```
INFO - Model directory: models
```

```
INFO -
```

```
You can now use this model with:
```

```
INFO - python src/predict.py document.pdf
```

```
INFO - =====
```

```
INFO - =====  
INFO - MODEL INSPECTION  
INFO - =====  
INFO - Model type: Pipeline  
INFO - Expected features: 28  
INFO - =====  
INFO -  
INFO - =====  
INFO - COMPATIBILITY CHECK  
INFO - =====  
INFO - Fitting preprocessing pipeline on datasets/PDF_All_features.csv...  
INFO - Fitting transformer on 32 features...  
INFO - ✓ Learned caps for 32 features  
INFO - ✓ Scaler fitted on 28 features  
INFO - ✓ Pipeline fitted and caps learned  
INFO - Our preprocessing outputs: 28 features  
INFO - Model expects: 28 features  
INFO - ✓ Feature count matches!  
INFO - =====  
INFO -  
INFO - =====  
INFO - MODEL TESTING  
INFO - =====  
INFO - Loading test data from datasets/PDF_All_features.csv...  
INFO - Loaded 19296 samples  
INFO - Applying our preprocessing...  
INFO - Making predictions...  
INFO - ✓ Predictions successful!  
INFO -
```

predict.py - Production Interface

```
python src/predict.py document.pdf
```

```
# 1. Load everything
```

```
model = load("models/model.joblib")
pipeline = PDFPreprocessingPipeline(
    caps=load("models/caps.json"),
    scaler=load("models/scaler.joblib")
)
```

```
# 2. Process PDF
```

```
raw_features = extractor.extract(pdf_path) # 41 features
X = pipeline.transform(raw_features) # 28 scaled features
```

```
# 3. Predict
```

```
probabilities = model.predict_proba(X)[0]
prediction = model.predict(X)[0]
```

Complete Flow

- Trained Model (from Model.ipynb)
- ↓
- import_model.py
- └ Fit pipeline on training data
- └ Save: model + caps + scaler
- ↓
- predict.py
- └ Load: model + caps + scaler
- └ Extract features from PDF
- └ Apply same preprocessing
- └ Predict: BENIGN/MALICIOUS

Result: Plug-and-play system for any trained model

Usage

- # One-time: Import trained model
 - python src/import_model.py model.joblib --test
-
- # Use: Predict on PDFs
 - python src/predict.py document.pdf
 - python src/predict.py document.pdf --json # API mode

Testing on Unseen Data

- Benign and Malicious PDFs were generated to test the models
- Benign PDFs were made using Microsoft Word
- Malicious PDFs were made using the DidierStevensSuite(<https://github.com/DidierStevens/DidierStevensSuite>)

file_path	label	malware_probability	prediction
---	---	---	---
str	str	f64	i64
pdf_files/test_output.pdf	Malware	0.999969	1
pdf_files/Shubham_Negi - Resum..	Benign	0.000071	0

Testing on Unseen Data

```
~/repositories/Trap4Phishing-ML pipeline *2 > python src/predict.py test_output.pdf
INFO - Loading model...
INFO - Using 28-feature (sparse removal) preprocessing mode
INFO - ✓ Model loaded
INFO - Analyzing: test_output.pdf
INFO - [1/3] Extracting features...
WARNING - PyPDF2 extraction failed for test_output.pdf: [Errno 22] Invalid argument
WARNING - PyMuPDF extraction failed for test_output.pdf: bad xref
INFO - [2/3] Applying preprocessing...
INFO - [3/3] Running model inference...
=====
⚠ PREDICTION: MALICIOUS
=====
File: test_output.pdf
Confidence: 95.3%

Probabilities:
  Benign: 4.7%
  Malicious: 95.3%
=====
~/repositories/Trap4Phishing-ML pipeline *2 ?1 > 
```

Testing on Unseen Data

```
INFO - 
~/repositories/Trap4Phishing-ML pipeline *2 !5 ?2 > python src/predict.py LearningAgreementSE.CNyandoro.pdf
INFO - Loading model...
INFO - Using 28-feature (sparse removal) preprocessing mode
INFO - ✓ Model loaded
INFO - Analyzing: LearningAgreementSE.CNyandoro.pdf
INFO - [1/3] Extracting features...
WARNING - PyMuPDF extraction failed for LearningAgreementSE.CNyandoro.pdf: bad xref
INFO - [2/3] Applying preprocessing...
INFO - [3/3] Running model inference...
=====
✓ PREDICTION: BENIGN
=====
File: LearningAgreementSE.CNyandoro.pdf
Confidence: 100.0%

Probabilities:
  Benign: 100.0%
  Malicious: 0.0%
=====
~/repositories/Trap4Phishing-ML pipeline *2 !5 ?2 >
```



Gracias por tu
atención y tiempo

Dataset Reference

Nejati, N. et al. "A Comprehensive Multi-Format Malicious Attachment Dataset for Email Threat Detection.“ Canadian Institute for Cybersecurity (CIC), University of New Brunswick, 2025.